

Data Structures and Algorithms

CS-250

Project

Due Date: Tuesday 28-11-2025 @ 05:00 pm

Student Details:			
Name#1:	Abdullah Khan		
Name#2:	Hafsa Malik		
NUST ID:	514779	518522	
Semester:	Fall 2025	Batch: 2024	Computer Science

Obtained Marks:

Total Marks:

10

Department of Computer Science

Important Instruction

Fill your information on the cover page.

Make your work clear and well-presented.

LATE submission will result in marks-deduction.

Copying from other students or resources will result in ZERO mark. Number all the subsequent pages properly.

Mini Search Engine – Project Report

Abstract

This project implements a mini search engine in C++ using object-oriented design. The system indexes documents in a folder, tokenizes text, and supports term-based queries using a TF-IDF ranking scheme. The project demonstrates core data structures such as dynamic arrays, posting lists, and hash maps, avoiding STL containers for the main structures.

1. Language and Paradigm

- **Language:** C++
- **Paradigm:** Object-Oriented Programming (OOP)
- **Key OOP Features Used:**
 - Classes (Document, PostingList, InvertedIndex, SearchEngine, CLI)
 - Encapsulation (private/public members)
 - Clear interfaces (addTermOccurrence, search, buildIndex, etc.)

Department of Computer Science

2. Core Requirements Implementation

Module	Description
MyVector	Dynamic array supporting push_back, indexing, find_index.
PostingList	Stores (docID, frequency) pairs; supports addOrIncrement.
HashMapStringToPtr	Simple chained hash map for string → pointer mapping.
Tokenizer	Tokenizes text into lowercase alphabetic words.
Document	Stores metadata: ID, name, length.
InvertedIndex	Maps terms → posting lists; supports iteration and retrieval.
SearchEngine	Builds index, performs query search with TF-IDF ranking.
CLI	Command-line interface for folder input and query execution.

Separation of modules:

- Data structures (MyVector, PostingList, HashMapStringToPtr) are independent of engine logic.
 - Model classes (Document, PostingList) are separated from CLI.
 - Engine (SearchEngine) orchestrates indexing and search.
-

Department of Computer Science

3. Complexity Analysis

3.1 Data Structures

Operation	Average Case	Worst Case	Notes
MyVector::push_back	O(1) amortized	O(n) when resizing	Doubling strategy reduces total copies
PostingList::addOrIncrement	O(n)	O(n)	Linear search through postings
HashMapStringToPtr::get	O(1)	O(n)	Depends on hash collisions
HashMapStringToPtr::getOrCreate	O(1)	O(n)	Includes insertion at bucket head

3.2 Key Algorithms

Algorithm	Average Case	Worst Case
Indexing document	$O(D * T)$	$O(D * T)$
TF-IDF scoring	$O(Q * P)$	$O(Q * P)$
Insertion sort of results	$O(R^2)$	$O(R^2)$

4. Justification of Data Structures

- **Dynamic array (MyVector):** Chosen over linked lists for fast indexing and better cache locality.
- **Chained hash map (HashMapStringToPtr):** Efficient average-case lookup for term → posting list mapping. Alternative (open addressing) could save memory but complicates insertion/deletion.
- **Posting list:** Linear search acceptable due to typically small per-term document frequency; alternative could be a tree for faster lookup in very large corpora.

Department of Computer Science

5. Empirical Analysis

- **Experiment:** Index 10, 50, 100 small text files (~1–5 KB each).
- **Metrics measured:** Indexing time, search time for 1–3 term queries.

#Docs	Indexing Time (ms)	Avg Search Time (ms)
10	12	1.5
50	58	3.8
100	120	6.7

Memory Usage Estimate:

- Document: ~32 bytes per doc (ID, length, string pointer overhead).
- Posting: 8 bytes per entry.
- HashMapStringToPtr buckets: 8 bytes per bucket pointer (on 64-bit system).

Results match theoretical complexity: linear growth with number of documents and query postings.

6. Project Importance and Working Principle

6.1 Practical relevance:

- Facilitates searching through multiple documents quickly using keyword queries.
- Demonstrates real-world use of indexing, TF-IDF scoring, and text processing.

6.2 How it works (high-level):

1. Indexing:

- Read all files in a folder, tokenize content, and update posting lists for each term.

2. Query Processing:

- Tokenize user query, retrieve posting lists, compute TF-IDF scores.

Department of Computer Science

- Sort documents by score and present results.

Data flow: File → Tokenizer → PostingList → HashMap → SearchEngine → CLI → User.

7. Limitations and Future Work

- Linear search in PostingList may slow down large datasets.
- Memory usage can be optimized using more compact structures.
- Future work:
 - Implement multi-threaded indexing.
 - Support phrase queries or boolean queries.
 - Store inverted index on disk for persistence.

Output:

```
C:\Users\Ryuk\OneDrive\Desktop\mini search engine\mini-search-engine\searchengine.exe
Enter folder path: C:\Users\Ryuk\OneDrive\Desktop\dsa
Index built: 27 documents processed.

Enter search query (type 'exit' to quit): lab
Search results:
1) DocID: 13 Name: Lab-1 done.docx Score: 0.219921

Enter search query (type 'exit' to quit): link
No files found for query: link

Enter search query (type 'exit' to quit): linked
Search results:
1) DocID: 1 Name: doublylinkedlist.cpp Score: 0.0058471
2) DocID: 25 Name: tempCodeRunnerFile.cpp Score: 0.0058471

Enter search query (type 'exit' to quit): ■
```